

Architecture of Enterprise Applications 13

MySQL Backup & Recovery

Haopeng Chen

***RE**liable, **IN**telligent and **SC**alable Systems Group (**REINS**)*

Shanghai Jiao Tong University

Shanghai, China

<http://reins.se.sjtu.edu.cn/~chenhp>

e-mail: chen-hp@sjtu.edu.cn

- Contents

- Backup and Recovery Types
- Database Backup Methods
- Example Backup and Recovery Strategy
- Using mysqldump for Backups
- Point-in-Time (Incremental) Recovery
- MyISAM Table Maintenance and Crash Recovery
- From: <https://dev.mysql.com/doc/refman/8.0/en/backup-and-recovery.html>



- Objectives

- 能够根据数据访问的具体场景，设计提高数据库访问性能和灾备能力的方案，包括集群部署和备份机制

- It is important to back up your databases
 - so that you can **recover** your data and be up and running again in case problems occur, such as **system crashes**, **hardware failures**, or **users deleting data by mistake**.
 - Backups are also essential as a **safeguard** before **upgrading** a MySQL installation, and they can be used to transfer a MySQL installation to another system or to set up replica servers.
- Several backup and recovery topics with which you should be familiar:
 - Types of backups: **Logical** versus **physical**, **full** versus **incremental**, and so forth.
 - Methods for **creating** backups.
 - **Recovery** methods, including point-in-time recovery.
 - Backup **scheduling**, **compression**, and **encryption**.
 - Table **maintenance**, to enable recovery of corrupt tables.

- Physical (Raw) Versus Logical Backups

- Physical backups consist of raw copies of the directories and files that store database contents.
 - This type of backup is suitable for **large, important databases** that need to be **recovered quickly** when problems occur.
- Logical backups save information represented as logical database structure ([CREATE DATABASE](#), [CREATE TABLE](#) statements) and content ([INSERT](#) statements or delimited-text files).
 - This type of backup is suitable for **smaller amounts of data** where you might edit the data values or table structure, or recreate the data **on a different machine architecture**.

- Physical backup methods have these characteristics:
 - The backup consists of **exact copies of database directories and files**. Typically this is a copy of all or part of the MySQL data directory.
 - Physical backup methods are **faster** than logical because they involve **only file copying without conversion**.
 - Output is **more compact** than for logical backup.
 - Because backup speed and compactness are important for **busy, important** databases, **the MySQL Enterprise Backup product performs physical backups**.
 - Backup and restore **granularity** ranges **from the level of the entire data directory down to the level of individual files**. This may or may not provide for table-level granularity, depending on storage engine.
 - For example, InnoDB tables can each be in **a separate file**, or share file storage with other InnoDB tables; each MyISAM table corresponds uniquely to **a set of files**.
 - In addition to databases, the backup can include any related files such as **log or configuration files**.

- Physical backup methods have these characteristics:
 - Data from **MEMORY** tables is tricky to back up this way
 - because their contents are **not** stored on disk.
 - Backups are **portable** only to other machines that have **identical or similar hardware characteristics**.
 - Backups can be performed while the MySQL server is **not running**.
 - If the server is running, it is necessary to perform appropriate **locking** so that the server does not change database contents during the backup.
 - MySQL Enterprise Backup does this locking **automatically** for tables that require it.
 - Physical backup tools include the **mysqlbackup** of MySQL Enterprise Backup for **InnoDB** or any other tables, or file system-level commands (such as **cp**, **scp**, **tar**, **rsync**) for **MyISAM** tables.
 - For restore:
 - MySQL Enterprise Backup restores **InnoDB** and other tables that it backed up.
 - **ndb_restore** restores **NDB** tables.
 - Files copied at the file system level can be copied back to their **original locations** with file system commands.

- Logical backup methods have these characteristics:
 - The backup is done by **querying** the MySQL server to **obtain database structure and content information**.
 - Backup is **slower** than physical methods because the server must **access** database information and **convert** it to logical format.
 - Output is **larger** than for physical backup, particularly when saved **in text format**.
 - Backup and restore **granularity** is available at the server level (all databases), database level (all tables in a particular database), or table level.
 - The backup does **not** include log or configuration files, or other database-related files that are **not** part of databases.
 - Backups stored in logical format are **machine independent and highly portable**.
 - Logical backups are performed with the MySQL server **running**.
 - Logical backup tools include the **mysqldump** program and the **SELECT ... INTO OUTFILE** statement. These work for any storage engine, even **MEMORY**.
 - To restore logical backups, **SQL-format dump files** can be processed using the **mysql** client. To load **delimited-text files**, use the **LOAD DATA** statement or the **mysqlimport** client.

- Online Versus Offline Backups

- **Online** backups take place while the MySQL server is **running** so that the database information can be obtained from the server.
- **Offline** backups take place while the server is **stopped**.
- This distinction can also be described as “**hot**” versus “**cold**” backups;
 - a “**warm**” backup is one where the server remains **running but locked against modifying data** while you access database files externally.
- **Online** backup methods have these characteristics:
 - The backup is **less intrusive to other clients**, which can connect to the MySQL server during the backup and may be able to access data depending on what operations they need to perform.
 - Care must be taken to **impose appropriate locking** so that data modifications do not take place that would compromise backup integrity. The MySQL Enterprise Backup product **does such locking automatically**.

- Online Versus Offline Backups

- Offline backup methods have these characteristics:

- Clients can be **affected adversely** because the server is **unavailable during backup**. For that reason, such backups are often taken from a **replica** that can be taken offline without harming availability.
 - The backup procedure is **simpler** because there is no possibility of interference from client activity.

- A similar distinction between **online** and **offline** applies for recovery operations, and similar characteristics apply.

- However, it is more likely for clients to be affected by online recovery than by online backup because **recovery requires stronger locking**.
 - During backup, clients might be able to read data while it is being backed up. Recovery modifies data and does not just read it, so clients must be **prevented from accessing data while it is being restored**.

- Local Versus Remote Backups

- A **local** backup is performed **on the same host** where the MySQL server runs,
 - whereas a **remote** backup is done from **a different host**. For some types of backups, the backup can be initiated from a remote host even if the output is written locally on the server host.
 - **mysqldump** can connect to local or remote servers.
 - For **SQL output** (CREATE and **INSERT** statements), local or remote dumps can be done and generate output **on the client**.
 - For **delimited-text output** (with the **--tab** option), data files are created **on the server host**.
 - **SELECT ... INTO OUTFILE** can be initiated from a local or remote client host, but the output file is created **on the server host**.
 - **Physical** backup methods typically are initiated **locally** on the MySQL server host so that the server can be taken **offline**, although the destination for copied files might be remote.

- Snapshot Backups

- Some file system implementations enable “**snapshots**” to be taken.
- These provide **logical copies** of the file system at a given point in time, **without requiring a physical copy of the entire file system**.
 - (For example, the implementation may use **copy-on-write** techniques so that only parts of the file system **modified after the snapshot** time need be copied.)
- MySQL itself **does not** provide the capability for taking file system snapshots.
 - It is available through **third-party** solutions such as Veritas, LVM, or ZFS.

- Full Versus Incremental Backups
 - A **full** backup includes all data managed by a MySQL server at a given point in time.
 - An **incremental** backup consists of the changes made to the data during a given time span (from one point in time to another).
 - MySQL has different ways to perform full backups.
 - Incremental backups are made possible by enabling the server's **binary log**, which the server uses **to record data changes**.

- Full Versus Point-in-Time (Incremental) Recovery
 - A **full** recovery restores **all** data from a full backup.
 - This restores the server instance **to the state that it had when the backup was made**.
 - If that state is not sufficiently current, a full recovery can be followed by recovery of incremental backups made since the full backup, to bring the server to a **more up-to-date state**.
 - Incremental recovery is recovery of **changes made during a given time span**.
 - This is also called **point-in-time recovery** because it makes a server's state current up to a given time. Point-in-time recovery is based on the **binary log** and typically follows a full recovery from the backup files that restores the server to its state when the backup was made.
 - Then the data changes written in the **binary log files** are applied as incremental recovery to **redo data modifications** and bring the server up to the desired point in time.

- Backup Scheduling, Compression, and Encryption
 - Backup scheduling is valuable for **automating backup procedures**.
 - Compression of backup output **reduces space requirements**, and
 - encryption of the output provides **better security against unauthorized access** of backed-up data.
 - MySQL itself **does not** provide these capabilities.
 - The MySQL Enterprise Backup product can compress InnoDB backups, and compression or encryption of backup output can be achieved using file system utilities. Other third-party solutions may be available.

- Making a Hot Backup with MySQL Enterprise Backup
 - Customers of **MySQL Enterprise Edition** can use the [MySQL Enterprise Backup](#) product to do [physical](#) backups of **entire** instances or **selected** databases, tables, or both.
 - This product includes features for [incremental](#) and [compressed](#) backups.
 - Backing up the **physical** database files makes restore much **faster** than **logical** techniques such as the **mysqldump** command.
 - InnoDB tables are copied using a [hot backup](#) mechanism.
 - (Ideally, the InnoDB tables should represent a substantial majority of the data.)
 - Tables from other storage engines are copied using a [warm backup](#) mechanism.

- Making Backups with `mysqldump`
 - The [`mysqldump`](#) program can make backups. It can back up all kinds of tables.
 - For InnoDB tables, it is possible to perform an online backup that takes no locks on tables using the [`--single-transaction`](#) option to [`mysqldump`](#).
- Making Backups by Copying Table Files
 - MyISAM tables can be backed up by **copying table files** (*.MYD, *.MYI files, and associated *.sdi files). To get a consistent backup, stop the server or lock and flush the relevant tables:
`FLUSH TABLES tbl_list WITH READ LOCK;`
 - You need only **a read lock**; this enables other clients to continue to query the tables while you are making a copy of the files in the database directory.
 - The flush is needed to ensure that the **all active index pages** are written to disk before you start the backup.
 - You can also create a binary backup simply by **copying the table files**, as long as the server isn't updating anything.
 - (But note that table file copying methods **do not work** if your database contains **InnoDB tables**. Also, even if the server is not actively updating data, InnoDB may still have modified data **cached in memory and not flushed to disk**.)

- Making Delimited-Text File Backups

- To create a **text file** containing a table's data, you can use [SELECT * INTO OUTFILE 'file name' FROM tbl name](#).
- The file is created on the MySQL server host, **not** the client host. For this statement, the output file **cannot** already exist because permitting files to be overwritten constitutes a security risk.
- This method works for any kind of data file, but saves **only table data**, **not** the **table structure**.
- Another way to create text data files (along with files containing [CREATE TABLE](#) statements for the backed up tables) is to use [mysqldump](#) with the [--tab](#) option.
- To reload a delimited-text data file, use [LOAD DATA](#) or [mysqlimport](#).

- Making Incremental Backups by Enabling the Binary Log
 - MySQL supports **incremental** backups using the **binary log**.
 - The binary log files provide you with the information you need to **replicate changes to the database** that are made subsequent to the point at which you performed a backup.
 - At the moment you want to make an **incremental** backup (containing all changes that happened since the last full or incremental backup), you should rotate the binary log by using [FLUSH LOGS](#).
 - The next time you do a full backup, you should also rotate the binary log using [FLUSH LOGS](#) or [mysqldump --flush-logs](#).

- Making Backups Using Replicas

- If you have performance problems with a server while making backups, one strategy that can help is to **set up replication** and **perform backups on the replica rather than on the source**.
- If you are backing up a replica, you should back up **its connection metadata repository** and applier metadata repository when you back up the replica's databases, regardless of the backup method you choose.
- This information is always needed to resume replication after you restore the replica's data.
 - If your replica is replicating [LOAD DATA](#) statements, you should also back up any **SQL_LOAD-*** files that exist in the directory that the replica uses for this purpose.
 - The replica needs these files to resume replication of any interrupted [LOAD DATA](#) operations.

- Recovering Corrupt Tables
 - If you have to restore **MyISAM** tables that have become corrupt, try to recover them using [REPAIR TABLE](#) or [myisamchk -r](#) first. That should work in **99.9%** of all cases.
- Making Backups Using a File System Snapshot
 - If you are using a **Veritas file system**, you can make a backup like this:
 - From a client program, execute [FLUSH TABLES WITH READ LOCK](#).
 - From another shell, execute mount vxfs snapshot.
 - From the first client, execute [UNLOCK TABLES](#).
 - Copy files from the snapshot.
 - Unmount the snapshot.
 - Similar snapshot capabilities may be available in other file systems, such as LVM or ZFS.

- Let's discuss a procedure for performing backups that enables you to recover data after several types of crashes:
 - Operating system crash
 - Power failure
 - File system crash
 - Hardware problem (hard drive, motherboard, and so forth)
- Assume that data is stored in the **InnoDB** storage engine, which has support for transactions and automatic crash recovery.
- Assume also that the MySQL server is **under load at the time of the crash**. If it were not, no recovery would ever be needed.

- For cases of **operating system crashes** or **power failures**, we can assume that **MySQL's disk data is available after a restart**.
 - The InnoDB data files might **not** contain **consistent** data due to the crash, but InnoDB **reads** its logs and **finds** in them the list of **pending committed** and **noncommitted transactions** that have not been flushed to the data files.
 - InnoDB **automatically rolls back** those transactions that were **not committed**, and **flushes** to its data files those that were **committed**.
- For the cases of **file system crashes** or **hardware problems**, we can assume that the MySQL disk data is **not** available after a restart.
 - This means that MySQL **fails to start** successfully because some blocks of disk data are **no longer readable**.
 - In this case, it is necessary to **reformat the disk**, **install a new one**, or otherwise **correct the underlying problem**.
 - Then it is necessary to recover our MySQL data **from backups**, which means that backups must already have been made. To make sure that is the case, design and implement a backup policy.

- Establishing a Backup Policy

- To be useful, backups must be scheduled regularly.
- Assume that we make **a full backup** of all our InnoDB tables in all databases using the following command on Sunday at 1 p.m., when load is low:

```
shell> mysqldump --all-databases --master-data --single-transaction > backup_sunday_1_PM.sql
```
- The resulting **.sql** file produced by [mysqldump](#) contains a set of SQL [INSERT](#) statements that can be used **to reload the dumped tables** at a later time.
- This backup operation acquires **a global read lock on all tables at the beginning of the dump.**
- Full backups are necessary, but it is **not** always convenient to create them.

- Establishing a Backup Policy

- To make incremental backups, we need to save the **incremental** changes.
- In MySQL, these changes are represented in the **binary log**,
 - so the MySQL server should always be started with the [--log-bin](#) option to enable that log. With binary logging enabled, the server writes **each data change** into a file while it updates data.

```
-rw-rw---- 1 guilhem guilhem 1277324 Nov 10 23:59 gbichot2-bin.000001
-rw-rw---- 1 guilhem guilhem      4 Nov 10 23:59 gbichot2-bin.000002
-rw-rw---- 1 guilhem guilhem    79 Nov 11 11:06 gbichot2-bin.000003
-rw-rw---- 1 guilhem guilhem   508 Nov 11 11:08 gbichot2-bin.000004
-rw-rw---- 1 guilhem guilhem 220047446 Nov 12 16:47 gbichot2-bin.000005
-rw-rw---- 1 guilhem guilhem  998412 Nov 14 10:08 gbichot2-bin.000006
-rw-rw---- 1 guilhem guilhem   361 Nov 14 10:07 gbichot2-bin.index
```

- The MySQL binary logs take up disk space. To free up space, purge them from time to time:

```
shell> mysqldump --single-transaction --flush-logs --master-data=2 \ -
--all-databases --delete-master-logs > backup_sunday_1_PM.sql
```


- Using Backups for Recovery

- Now, suppose that we have a **catastrophic unexpected exit** on Wednesday at 8 a.m. that requires recovery from backups.
- To recover, first we restore **the last full backup** we have (the one from Sunday 1 p.m.):

```
shell> mysql < backup_sunday_1_PM.sql
```
- At this point, the data is restored to its state as of Sunday 1 p.m.
- To restore the changes made since then, we must use the **incremental** backups; that is, the **gbichot2-bin.000007** and **gbichot2-bin.000008** binary log files.
- Fetch the files if necessary from where they were backed up, and then process their contents like this:

```
shell> mysqlbinlog gbichot2-bin.000007 gbichot2-bin.000008 | mysql
```

- Using Backups for Recovery

- We now have recovered the data to its state as of **Tuesday 1 p.m.**, but still are missing the changes from that date to the date of the crash.
- To **not** lose them, we would have needed to have the MySQL server store its MySQL binary logs into **a safe location** (RAID disks, SAN, ...) **different from the place where it stores its data files**, so that these logs were **not** on the destroyed disk.
- (That is, we can start the server with a **--log-bin** option that specifies a location on a different physical device from the one on which the data directory resides.)
- If we had done this, we would have the **gbichot2-bin.000009** file (and any subsequent files) at hand, and we could apply them using **mysqlbinlog** and **mysql** to restore the most recent data changes with no loss up to the moment of the crash:

```
shell> mysqlbinlog gbichot2-bin.000009 ... | mysql
```

- Backup Strategy Summary

- In case of an **operating system crash** or **power failure**, InnoDB **itself** does all the job of recovering data. But to make sure that you can sleep well, observe the following guidelines:
- Always tun the MySQL server with **binary logging enabled** (that is the default setting for MySQL 8.0).
 - If you have such safe media, this technique can **also be good for disk load balancing** (which results in a performance improvement).
- Make **periodic full backups**, using the [mysqldump](#) command that makes an online, nonblocking backup.
- Make **periodic incremental backups** by flushing the logs with [FLUSH LOGS](#) or [mysqladmin flush-logs](#).

- Consider using the [MySQL Shell dump utilities](#), which provide
 - parallel dumping with **multiple threads**, **file compression**, and **progress information display**, as well as **cloud features** such as Oracle Cloud Infrastructure Object Storage streaming, and **MySQL Database Service compatibility checks** and **modifications**.
 - Dumps can be easily imported into a MySQL Server instance or a MySQL Database Service DB System using the [MySQL Shell load dump utilities](#).
- A dump file can be used in several ways:
 - As a **backup** to enable data recovery in case of data loss.
 - As a source of data for **setting up replicas**.
 - As a source of data for **experimentation**:
 - To make a copy of a database that you can use **without changing the original data**.
 - To test **potential upgrade incompatibilities**.

- [mysqldump](#) produces two types of output, depending on whether the [--tab](#) option is given:
 - Without [--tab](#), [mysqldump](#) writes SQL statements to the standard output.
 - This output consists of **CREATE** statements to create dumped objects (databases, tables, stored routines, and so forth), and **INSERT** statements to load data into tables.
 - The output can be saved in **a file** and reloaded later using [mysql](#) to recreate the dumped objects. Options are available to modify the format of the SQL statements, and to control which objects are dumped.
 - With [--tab](#), [mysqldump](#) produces two output files for each dumped table.
 - The server writes one file as **tab-delimited text, one line per table row**. This file is named **tbl_name.txt** in the output directory.
 - The server also sends a [CREATE TABLE](#) statement for the table to [mysqldump](#), which writes it as a file named **tbl_name.sql** in the output directory.

- Dumping Data in SQL Format with **mysqldump**

- By default, **mysqldump** writes information as SQL statements to the standard output. You can save the output in a file:

```
shell> mysqldump [arguments] > file_name
```

- To dump **all** databases, invoke **mysqldump** with the **--all-databases** option:

```
shell> mysqldump --all-databases > dump.sql
```

- To dump **only specific databases**, name them on the command line and use the **--databases** option:

```
shell> mysqldump --databases db1 db2 db3 > dump.sql
```

- To dump **a single database**, name it on the command line:

```
shell> mysqldump --databases test > dump.sql
```

- In the **single-database** case, it is permissible to omit the **--databases** option:

```
shell> mysqldump test > dump.sql
```

- To dump **only specific tables from a database**, name them on the command line following the database name:

```
shell> mysqldump test t1 t3 t7 > dump.sql
```

- Reloading SQL-Format Backups

- If the dump file was created by [mysqldump](#) with the [--all-databases](#) or [--databases](#) option, it contains [CREATE DATABASE](#) and [USE](#) statements and it is **not** necessary to specify a default database into which to load the data:

```
shell> mysql < dump.sql
```

- Alternatively, from within [mysql](#), use a source command:

```
mysql> source dump.sql
```

- If the file **is a single-database dump not** containing [CREATE DATABASE](#) and [USE](#) statements, create the database first (if necessary):

```
shell> mysqladmin create db1
```

- Then **specify the database name** when you load the dump file:

```
shell> mysql db1 < dump.sql
```

- Alternatively, from within [mysql](#), create the database, select it as the default database, and load the dump file:

```
mysql> CREATE DATABASE IF NOT EXISTS db1; mysql> USE db1; mysql> source  
dump.sql
```

- Dumping Data in Delimited-Text Format with mysqldump
 - If you invoke [mysqldump](#) with the [--tab=dir_name](#) option, it uses *dir_name* as the output directory and dumps tables individually in that directory using **two files for each table**. The table name is the base name for these files.
 - For a table named **t1**, the files are named **t1.sql** and **t1.txt**. The **.sql** file contains a [CREATE TABLE](#) statement for the table. The **.txt** file contains the table data, one line per table row.
 - The following command dumps the contents of the db1 database to files in the /tmp database:

```
shell> mysqldump --tab=/tmp db1
```


- Reloading Delimited-Text Format Backups

- To reload a table, first change location into the output directory. Then process the **.sql** file with [mysql](#) to create an empty table and process the **.txt** file to load the data into the table:

```
shell> mysql db1 < t1.sql
```

```
shell> mysqlimport db1 t1.txt
```

- An alternative to using [mysqlimport](#) to load the data file is to use the [LOAD DATA](#) statement from within the [mysql](#) client:

```
mysql> USE db1;
```

```
mysql> LOAD DATA INFILE 't1.txt' INTO TABLE t1;
```

- If you used any data-formatting options with [mysqldump](#) when you initially dumped the table, you must use the same options with [mysqlimport](#) or [LOAD DATA](#) to ensure proper interpretation of the data file contents:

```
shell> mysqlimport --fields-terminated-by=, --fields-enclosed-by='"' --lines-terminated-by=0x0d0a db1 t1.txt
```

- Or:

```
mysql> USE db1;
```

```
mysql> LOAD DATA INFILE 't1.txt' INTO TABLE t1 FIELDS TERMINATED BY ',' FIELDS ENCLOSED BY '"' LINES TERMINATED BY '\r\n';
```

- How to make a copy a database
- How to copy a database from one server to another
- How to dump stored programs (stored procedures and functions, triggers, and events)
- How to dump definitions and data separately

- Making a Copy of a Database

- shell> `mysqldump db1 > dump.sql`
 - shell> `mysqladmin create db2`
 - shell> `mysql db2 < dump.sql`

- Do not use `--databases` on the `mysqldump` command line because that causes **USE db1** to be included in the dump file, which overrides the effect of naming **db2** on the `mysql` command line.

- Copy a Database from one Server to Another

- On Server 1:

- ```
shell> mysqldump --databases db1 > dump.sql
```

- Copy the dump file from Server 1 to Server 2.

- On Server 2:

- ```
shell> mysql < dump.sql
```

- Alternatively, you can omit [--databases](#) from the [mysqldump](#) command.

- On Server 1:

- ```
shell> mysqldump db1 > dump.sql
```

- On Server 2:

- ```
shell> mysqladmin create db1
```

- ```
shell> mysql db1 < dump.sql
```

- Dumping Stored Programs

- Several options control how [mysqldump](#) handles stored programs (**stored procedures and functions, triggers, and events**):
  - [--events](#): Dump Event Scheduler events
  - [--routines](#): Dump stored procedures and functions
  - [--triggers](#): Dump triggers for tables
- The [--triggers](#) option is **enabled by default** so that when tables are dumped, they are accompanied by any triggers they have.
  - **The other options are disabled by default** and must be specified explicitly to dump the corresponding objects.
  - To disable any of these options explicitly, use its skip form: [--skip-events](#), [--skip-routines](#), or [--skip-triggers](#).

- Dumping Table Definitions and Content Separately
  - The `--no-data` option tells `mysqldump` **not** to dump table data, resulting in the dump file containing **only statements to create the tables**.
    - Conversely, the `--no-create-info` option tells `mysqldump` to **suppress CREATE** statements from the output, so that the dump file contains **only table data**.
  - For example, to dump table **definitions and data separately** for the test database, use these commands:

```
shell> mysqldump --no-data test > dump-defs.sql
shell> mysqldump --no-create-info test > dump-data.sql
```
  - For a **definition-only dump**, add the `--routines` and `--events` options to also include **stored routine** and **event definitions**:

```
shell> mysqldump --no-data --routines --events test > dump-defs.sql
```

- Using mysqldump to Test for Upgrade Incompatibilities

- When contemplating a MySQL **upgrade**, it is prudent to install the newer version separately from your current production version.
  - Then you can dump the database and database object definitions from the production server and **load** them into the new server to verify that they are handled properly. (This is also useful for testing downgrades.)
- On the production server:

```
shell> mysqldump --all-databases --no-data --routines --events > dump-
defs.sql
```
- On the upgraded server:

```
shell> mysql < dump-defs.sql
```
- Because the dump file **does not contain table data**, it can be processed **quickly**.

- Using mysqldump to Test for Upgrade Incompatibilities
  - This enables you to spot potential **incompatibilities** without waiting for lengthy data-loading operations. Look for warnings or errors while the dump file is being processed.
  - After you have verified that the definitions are handled properly, dump the data and try to load it into the upgraded server.
  - On the production server:

```
shell> mysqldump --all-databases --no-create-info > dump-data.sql
```
  - On the upgraded server:

```
shell> mysql < dump-data.sql
```
  - Now check the table contents and run some test queries.



- Point-in-time recovery refers to
  - recovery of data changes up to a given point in time.
  - Typically, this type of recovery is performed **after restoring a full backup** that brings the server to its state as of the time the backup was made.
  - Point-in-time recovery then brings the server up to **date incrementally from the time of the full backup to a more recent time.**

- Point-in-Time Recovery Using Binary Log

- To restore data from the binary log, you must know the name and location of the current binary log files.

```
mysql> SHOW BINARY LOGS;
```

- To determine the name of the **current binary log file**, issue the following statement:

```
mysql> SHOW MASTER STATUS;
```

- To apply events from the binary log, process [mysqlbinlog](#) output using the [mysql](#) client:

```
shell> mysqlbinlog binlog_files | mysql -u root -p
```

- If binary log files have been encrypted, which can be done from MySQL 8.0.14 onwards, [mysqlbinlog](#) **cannot** read them directly as in the above example, but can read them from the server using the [--read-from-remote-server](#) (-R) option. For example:

```
shell> mysqlbinlog --read-from-remote-server --host=host_name --
port=3306 --user=root --password --ssl-mode=required binlog_files |
mysql -u root -p
```

- Point-in-Time Recovery Using Binary Log

- To **view** events from the log, send [mysqlbinlog](#) output into a paging program:

```
shell> mysqlbinlog binlog_files | more
```

- Alternatively, save the output in a file and view the file in a text editor:

```
shell> mysqlbinlog binlog_files > tmpfile
```

```
shell> ... edit tmpfile ...
```

- **Saving** the output in a file is useful as a preliminary to executing the log contents **with certain events removed**, such as an accidental [DROP TABLE](#).
- You can delete from the file any statements **not** to be executed before executing its contents.
- After editing the file, apply the contents as follows:

```
shell> mysql -u root -p < tmpfile
```

- Point-in-Time Recovery Using Binary Log

- If you have **more than one binary log** to apply on the MySQL server, the safe method is to process them all using a single connection to the server.

- Here is an example that demonstrates what may be *unsafe*:

```
shell> mysqlbinlog binlog.000001 | mysql -u root -p # DANGER!!
```

```
shell> mysqlbinlog binlog.000002 | mysql -u root -p # DANGER!!
```

- Processing binary logs this way using **different connections** to the server causes problems if the **first** log file contains a CREATE TEMPORARY TABLE statement and the **second** log contains a statement that uses the temporary table.
  - When the **first mysql** process terminates, the server **drops the temporary table**.
  - When the **second mysql** process attempts to use the table, the server reports “**unknown table**.”

- Point-in-Time Recovery Using Binary Log

- To avoid problems like this, use a *single* connection to apply the contents of all binary log files that you want to process.

- Here is one way to do so:

```
shell> mysqlbinlog binlog.000001 binlog.000002 | mysql -u root -p
```

- Another approach is to write the whole log to a single file and then process the file:

```
shell> mysqlbinlog binlog.000001 > /tmp/statements.sql
```

```
shell> mysqlbinlog binlog.000002 >> /tmp/statements.sql
```

```
shell> mysql -u root -p -e "source /tmp/statements.sql"
```

- When writing to a dump file while reading back from a binary log containing GTIDs ([“Replication with Global Transaction Identifiers”](#)), use the [--skip-gtids](#) option with [mysqlbinlog](#), like this:

```
shell> mysqlbinlog --skip-gtids binlog.000001 > /tmp/dump.sql
```

```
shell> mysqlbinlog --skip-gtids binlog.000002 >> /tmp/dump.sql
```

```
shell> mysql -u root -p -e "source /tmp/dump.sql"
```

- Point-in-Time Recovery Using Event Positions
  - As an example, suppose that around **20:06:00 on March 11, 2020**, an SQL statement was executed that **deleted a table**.
    - You can perform a point-in-time recovery to **restore the server** up to its state **right before the table deletion**.
  - These are some sample steps to achieve that:
    - Restore the **last full backup** created before the point-in-time of interest (call it  $t_p$ , which is 20:06:00 on March 11, 2020 in our example). When finished, note the binary log position up to which you have restored the server for later use, and **restart the server**.
    - Find the **precise binary log event position** corresponding to the point in time up to which you want to restore your database.
      - In our example, given that we know the rough time where the table deletion took place ( $t_p$ ), we can find the log position by checking the log contents around that time using the [mysqlbinlog](#) utility.
      - Use the [--start-datetime](#) and [--stop-datetime](#) options to **specify a short time period around  $t_p$** , and then look for the event in the output.

- Point-in-Time Recovery Using Event Positions

```
shell> mysqlbinlog --start-datetime="2020-03-11 20:05:00" \
--stop-datetime="2020-03-11 20:08:00" --verbose \
/var/lib/mysql/bin.123456 | grep -C 15 "DROP TABLE"

/*!80014 SET @@session.original_server_version=80019/*!*/;
/*!80014 SET @@session.immediate_server_version=80019/*!*/;
SET @@SESSION.GTID_NEXT= 'ANONYMOUS'/*!*/;
at 232
#200311 20:06:20 server id 1 end_log_pos 355 CRC32 0x2fc1e5ea Query thread_id=16 exec_time=0 error_code=0
SET TIMESTAMP=1583971580/*!*/;
SET @@session.pseudo_thread_id=16/*!*/;
SET @@session.foreign_key_checks=1, @@session.sql_auto_is_null=0, @@session.unique_checks=1, @@session.autocommit=1/*!*/;
SET @@session.sql_mode=1168113696/*!*/;
SET @@session.auto_increment_increment=1, @@session.auto_increment_offset=1/*!*/;
/*!!\C utf8mb4 *//*!*/;
SET @@session.character_set_client=255,@@session.collation_connection=255,@@session.collation_server=255/*!*/;
SET @@session.lc_time_names=0/*!*/;
SET @@session.collation_database=DEFAULT/*!*/;
/*!80011 SET @@session.default_collation_for_utf8mb4=255/*!*/;
DROP TABLE `pets`.`cats` /* generated by server */
/*!*/;
at 355
#200311 20:07:48 server id 1 end_log_pos 434 CRC32 0x123d65df Anonymous_GTID last_committed=1 sequence_number=2
original_commit_timestamp=1583971668462467 (2020-03-11 20:07:48.462467 EDT)
immediate_commit_timestamp=1583971668462467 (2020-03-11 20:07:48.462467 EDT)
/*!80001 SET @@session.original_commit_timestamp=1583971668462467/*!*/;
/*!80014 SET @@session.original_server_version=80019/*!*/;
/*!80014 SET @@session.immediate_server_version=80019/*!*/;
SET @@SESSION.GTID_NEXT= 'ANONYMOUS'/*!*/;
at 434
#200311 20:07:48 server id 1 end_log_pos 828 CRC32 0x57fac9ac Query thread_id=16 exec_time=0 error_code=0 Xid =
use `pets`/*!*/;
SET TIMESTAMP=1583971668/*!*/;
/*!80013 SET @@session.sql_require_primary_key=0/*!*/;
CREATE TABLE dogs
```

- Point-in-Time Recovery Using Event Positions

- Apply the events in binary log file to the server, starting with the log position you found in step 1 (assume it is 155) and ending at the position you have found in step 2 that is *before* your point-in-time of interest (which is 232):

```
shell> mysqlbinlog --start-position=155 --stop-position=232
/var/lib/mysql/bin.123456 \ | mysql -u root -p
```

- Your database has now been restored to the point-in-time of interest,  $t_p$ , **right before the table `pets.cats` was dropped.**
- Beyond the point-in-time recovery that has been finished, if you also want to reexecute all the statements **after** your point-in-time of interest, use [`mysqlbinlog`](#) again to apply all the events **after  $t_p$**  to the server.
- We noted in step 2 that after the statement we wanted to skip, the log is at position 355; we can use it for the [`--start-position`](#) option, so that any statements after the position are included:

```
shell> mysqlbinlog --start-position=355 /var/lib/mysql/bin.123456 \ | mysql
-u root -p
```



- MyISAM Table Maintenance and Crash Recovery
  - You can use [myisamchk](#) to **check**, **repair**, or **optimize** database tables.
  - MyISAM table maintenance can also be done using the SQL statements that perform operations similar to what [myisamchk](#) can do:
    - To check MyISAM tables, use [CHECK TABLE](#).
    - To repair MyISAM tables, use [REPAIR TABLE](#).
    - To optimize MyISAM tables, use [OPTIMIZE TABLE](#).
    - To analyze MyISAM tables, use [ANALYZE TABLE](#).
  - One advantage of these statements over [myisamchk](#) is that **the server does all the work**.
    - With [myisamchk](#), you must make sure that the server **does not** use the tables at the same time so that there is no unwanted interaction between [myisamchk](#) and the server.

- Using myisamchk for Crash Recovery

- If you run [mysqld](#) with **external locking disabled** (which is the **default**), you **cannot** reliably use [myisamchk](#) to check a table when [mysqld](#) is using the **same** table.
  - If you can be **certain** that **no one** can access the tables using [mysqld](#) while you run [myisamchk](#), you only have to execute [mysqladmin flush-tables](#) before you start checking the tables.
  - If you **cannot guarantee** this, you must **stop** [mysqld](#) while you check the tables.
  - If you run [myisamchk](#) to check tables that [mysqld](#) is updating at the same time, you may get **a warning that a table is corrupt even when it is not**.
- If the server is run with **external locking enabled**, you can use [myisamchk](#) to check tables **at any time**.
  - In this case, if the server tries to update a table that [myisamchk](#) is using, the server **waits for** [myisamchk](#) to finish before it continues.

- Using `myisamchk` for Crash Recovery

- If you use `myisamchk` to **repair** or **optimize** tables, you **must** always ensure that the `mysqld` server is **not** using the table (this also applies if external locking is disabled).
  - If you **do not stop** `mysqld`, you should at least do a `mysqladmin flush-tables` before you run `myisamchk`. Your tables **may become corrupted** if the server and `myisamchk` access the tables simultaneously.
- `myisamchk` works by **creating a copy of the .MYD data file row by row**.
  - It ends the repair stage by **removing the old .MYD file** and renaming the new file to the original file name.
  - If you use `--quick`, `myisamchk` does not create a temporary .MYD file, but instead assumes that the .MYD file is correct and generates only a new index file without touching the .MYD file.
  - This is **safe**, because `myisamchk` automatically detects whether the .MYD file is corrupt and aborts the repair if it is.

- How to Check MyISAM Tables for Errors

- [myisamchk tbl name](#)

- This finds **99.99% of all errors**. What it **cannot** find is corruption that involves **only** the data file (which is very unusual). If you want to check a table, you should normally run [myisamchk](#) without options or with the -s (silent) option.

- [myisamchk -m tbl name](#)

- This finds **99.999% of all errors**. It **first** checks **all index entries** for errors and **then** reads through **all rows**. It calculates a checksum for all key values in the rows and verifies that the checksum matches the checksum for the keys in the index tree.

- [myisamchk -e tbl name](#)

- This does a **complete and thorough check of all data** (-e means “extended check”). It does a check-read of every key for each row to verify that they indeed point to the correct row. This may take **a long time** for a large table that has many indexes. Normally, [myisamchk](#) stops after the first error it finds. If you want to obtain more information, you can add the -v (verbose) option. This causes [myisamchk](#) to keep going, up through a maximum of 20 errors.

- [myisamchk -e -i tbl name](#)

- This is like the previous command, but the -i option tells [myisamchk](#) to **print additional statistical information**.

- In most cases, a simple [myisamchk](#) command with **no arguments** other than the table name is sufficient to check a table.

- How to Repair MyISAM Tables

- Symptoms of corrupted tables include queries that abort unexpectedly and observable errors such as these:
  - Can't find file *tbl\_name.MYI* (Errcode: *nnn*)
  - Unexpected end of file
  - Record file is crashed
  - Got error *nnn* from table handler
- To get more information about the error, run  `perror nnn`, where *nnn* is the error number:

```
shell> perror 126 127 132 134 135 136 141 144 145
MySQL error code 126 = Index file is crashed
MySQL error code 127 = Record-file is crashed
MySQL error code 132 = Old database file
MySQL error code 134 = Record was already deleted (or record file crashed)
MySQL error code 135 = No more room in record file
MySQL error code 136 = No more room in index file
MySQL error code 141 = Duplicate unique key or constraint on write or update
MySQL error code 144 = Table is crashed and last repair failed
MySQL error code 145 = Table was marked as crashed and should be repaired
```

- How to Repair MyISAM Tables
- ***Stage 1: Checking your tables***
  - Run [`myisamchk \*.MYI`](#) or [`myisamchk -e \*.MYI`](#) if you have more time. Use the -s (silent) option to suppress unnecessary information.
  - If the [`mysqld`](#) server is **stopped**, you should use the [`--update-state`](#) option to tell [`myisamchk`](#) to mark the table as “checked.”
  - You have to **repair** only those tables for which [`myisamchk`](#) announces an **error**. For such tables, proceed to Stage 2.
  - If you get **unexpected errors when checking** (such as out of memory errors), or if [`myisamchk`](#) **crashes**, go to Stage 3.

- How to Repair MyISAM Tables
- ***Stage 2: Easy safe repair***
  - First, try `myisamchk -r -q tbl name` (-r -q means “**quick recovery mode**”). This attempts to **repair the index file without touching the data file**. If the data file contains everything that it should and the delete links point at the correct locations within the data file, this should work, and the table is fixed. Start repairing the next table. Otherwise, use the following procedure:
    - Make a **backup** of the data file before continuing.
    - Use `myisamchk -r tbl name` (-r means “recovery mode”). This **removes incorrect rows and deleted rows from the data file and reconstructs the index file**.
    - If the preceding step **fails**, use `myisamchk --safe-recover tbl name`. Safe recovery mode uses **an old recovery method** that handles a few cases that regular recovery mode does not (**but is slower**).
  - If you get **unexpected errors when repairing** (such as out of memory errors), or if `myisamchk` crashes, go to Stage 3.

- How to Repair MyISAM Tables
- ***Stage 3: Difficult repair***
  - You should reach this stage **only if the first 16KB block in the index file is destroyed or contains incorrect information**, or if **the index file is missing**. In this case, it is necessary to create a new index file. Do so as follows:
    - Move the data file to a safe place.
    - Use the table description file to create new (empty) data and index files:

```
shell> mysql db_name
mysql> SET autocommit=1;
mysql> TRUNCATE TABLE tbl_name;
mysql> quit
```
    - Copy the old data file back onto the newly created data file. (Do not just move the old file back onto the new file. You want to retain a copy in case something goes wrong.)
  - Go back to Stage 2. [myisamchk -r -q](#) should work. (This should not be an endless loop.)



- 请你根据上课内容，针对你在E-BookStore项目中的数据库设计，详细回答下列问题：
  1. 请你详细描述如何通过全量备份和增量备份来实现系统状态恢复。(2分)
  2. 请你根据MySQL缓存的工作原理，描述预取机制的优点。(1分)
  - 请提交包含上述问题答案的文档，答案应该结合你的E-BookStore的具体设计来阐述，不要过于泛化。
- 评分标准：
  - 分值如问题描述，答案不唯一，只要你的说理合理即可视为正确。



Thank You!