

Architecture of Enterprise Applications 12

MySQL Optimization II

Haopeng Chen

***RE**liable, **IN**telligent and **SC**alable Systems Group (**REINS**)*

Shanghai Jiao Tong University

Shanghai, China

<http://reins.se.sjtu.edu.cn/~chenhp>

e-mail: chen-hp@sjtu.edu.cn

- Contents

- Optimizing for InnoDB Tables
- Optimizing for MyISAM Tables
- Optimizing for MEMORY Tables
- From: <https://dev.mysql.com/doc/refman/8.0/en/optimization.html>



- Objectives

- 能够根据数据访问的具体场景，设计数据库性能调优方案，包括索引优化、缓存设置和参数调优等

- Optimizing Storage Layout for InnoDB Tables
 - Once your data reaches a **stable size**, or a growing table has increased by tens or some hundreds of megabytes, consider using the **OPTIMIZE TABLE** statement to reorganize the table and compact any wasted space.
 - **OPTIMIZE TABLE** copies the data part of the table and rebuilds the indexes. The benefits come from improved packing of data within indexes, and reduced fragmentation within the tablespaces and on disk.
 - In InnoDB, having a **long PRIMARY KEY** (either a single column with a lengthy value, or several columns that form a long composite value) wastes a lot of disk space.
 - Create an **AUTO_INCREMENT** column as the primary key if your primary key is long, or index a prefix of a long VARCHAR column instead of the entire column.
 - Use the **VARCHAR** data type instead of **CHAR** to store variable-length strings or for columns with **many NULL values**.
 - For tables that are **big**, or contain **lots of repetitive text or numeric data**, consider using **COMPRESSED** row format.

- Optimizing InnoDB Transaction Management
 - The default MySQL setting **AUTOCOMMIT=1** can impose performance limitations on a busy database server.
 - Where practical, **wrap several related data change operations into a single transaction**, by issuing **SET AUTOCOMMIT=0** or a **START TRANSACTION** statement, followed by a **COMMIT** statement after making all the changes.
 - Alternatively, for transactions that consist only of a single **SELECT** statement, turning on AUTOCOMMIT helps InnoDB to recognize read-only transactions and optimize them.
 - **Avoid** performing rollbacks after **inserting, updating, or deleting huge numbers of rows**.
 - If a big transaction is slowing down server performance, rolling it back can make the problem worse, potentially taking several times as long to perform as the original data change operations.
 - Killing the database process does not help, because the rollback starts again on server startup.

- Optimizing InnoDB Transaction Management
 - To minimize the chance of this issue occurring:
 - Increase the size of the [buffer pool](#) so that all the data change changes can be cached rather than immediately written to disk.
 - Set [innodb change buffering=all](#) so that update and delete operations are buffered in addition to inserts.
 - Consider issuing **COMMIT** statements periodically during the big data change operation, possibly breaking a single delete or update into multiple statements that operate on smaller numbers of rows.
 - **A long-running transaction** can prevent InnoDB from purging data that was changed by a different transaction.
 - When rows are modified or deleted within a long-running transaction, other transactions using the [READ COMMITTED](#) and [REPEATABLE READ](#) isolation levels have to do more work to reconstruct the older data if they read those same rows.
 - When a long-running transaction modifies a table, queries against that table from other transactions do not make use of the [covering index](#) technique.

- Optimizing InnoDB Read-Only Transactions
 - InnoDB can avoid the overhead associated with setting up the [transaction ID](#) (TRX_ID field) for transactions that are known to be read-only.
 - InnoDB detects read-only transactions when:
 - The transaction is started with the [START TRANSACTION READ ONLY](#) statement.
 - The [autocommit](#) setting is turned on, so that the transaction is guaranteed to be a single statement, and the single statement making up the transaction is a “non-locking” [SELECT](#) statement.
 - The transaction is started without the **READ ONLY** option, but **no updates or statements that explicitly lock rows** have been executed yet.
 - Thus, for a read-intensive application such as a report generator, you can tune a sequence of InnoDB queries
 - by grouping them inside [START TRANSACTION READ ONLY](#) and [COMMIT](#), or
 - by turning on the [autocommit](#) setting before running the **SELECT** statements, or
 - simply by avoiding any data change statements interspersed with the queries.

- Optimizing InnoDB Redo Logging

- Make your redo log files big, even as big as the [buffer pool](#).
 - The size and number of redo log files are configured using the [innodb log file size](#) and [innodb log files in group](#) configuration options.
- Consider increasing the size of the [log buffer](#).
 - Log buffer size is configured using the [innodb log buffer size](#) configuration option, which can be configured dynamically in MySQL 8.0.
- Configure the [innodb log write ahead size](#) configuration option to avoid “read-on-write”.
 - Set [innodb log write ahead size](#) to match the operating system or file system cache block size.
 - Valid values for [innodb log write ahead size](#) are multiples of the InnoDB log file block size (2^n).
- MySQL 8.0.11 introduced dedicated log writer threads for writing redo log records from the log buffer to the system buffers and flushing the system buffers to the redo log files
 - As of MySQL 8.0.22, you can enable or disable log writer threads using the [innodb log writer threads](#) variable.
- Optimize the use of spin delay by user threads waiting for flushed redo.
 - [innodb log wait for flush spin hwm](#): Defines the maximum average log flush time beyond which user threads no longer spin while waiting for flushed redo. The default value is 400 microseconds.
 - [innodb log spin cpu abs lwm](#): Defines the minimum amount of CPU usage below which user threads no longer spin while waiting for flushed redo. The value is expressed as a sum of CPU core usage.
 - [innodb log spin cpu pct hwm](#): Defines the maximum amount of CPU usage above which user threads no longer spin while waiting for flushed redo.
 - The [innodb log spin cpu pct hwm](#) configuration option respects processor affinity. For example, if a server has 48 cores but the [mysqld](#) process is pinned to only four CPU cores, the other 44 CPU cores are ignored.

- Bulk Data Loading for InnoDB Tables

- When importing data into InnoDB, turn off **autocommit** mode, because it performs **a log flush to disk for every insert**.

```
SET autocommit=0;
```

```
... SQL import statements ...
```

```
COMMIT;
```

- If you have **UNIQUE** constraints on secondary keys, you can speed up table imports by **temporarily turning off the uniqueness checks** during the import session:

```
SET unique_checks=0;
```

```
... SQL import statements ...
```

```
SET unique_checks=1;
```

For big tables, this saves a lot of disk I/O because InnoDB can use its change buffer to write secondary index records in a batch. **Be certain that the data contains no duplicate keys.**

- Bulk Data Loading for InnoDB Tables

- If you have **FOREIGN KEY** constraints in your tables, you can speed up table imports by **turning off the foreign key checks** for the duration of the import session:

```
SET foreign_key_checks=0;
```

```
... SQL import statements ...
```

```
SET foreign_key_checks=1;
```

For big tables, this can save a lot of disk I/O.

- Use the **multiple-row INSERT** syntax to reduce communication overhead between the client and the server if you need to insert many rows:

```
INSERT INTO yourtable VALUES (1,2), (5,5), ...;
```

This tip is valid for inserts into any table, not just InnoDB tables.

- When doing bulk inserts into tables with auto-increment columns, set [innodb autoinc lock mode](#) to 2 (interleaved) instead of 1 (consecutive).

- Optimizing InnoDB Queries

- Because each InnoDB table has a [primary key](#) (whether you request one or not), specify a set of primary key columns for each table, columns that are used in the most important and time-critical queries.
- Do **not** specify too many or too long columns in the primary key, because these column values are duplicated in each secondary index.
- Do **not** create a separate [secondary index](#) for each column, because each query can only make use of one index.
 - If you have many queries for the same table, testing different combinations of columns, try to create a small number of [concatenated indexes](#) rather than a large number of single-column indexes.
- If an indexed column **cannot** contain any **NULL** values, declare it as **NOT NULL** when you create the table.
- You can optimize single-query transactions for InnoDB tables, [Optimizing InnoDB Read-Only Transactions](#)

- Optimizing InnoDB DDL Operations

- Use [TRUNCATE TABLE](#) to **empty** a table, **not** DELETE FROM *tbl_name*.
 - Foreign key constraints can make a TRUNCATE statement work like a regular **DELETE** statement, in which case a sequence of commands like [DROP TABLE](#) and [CREATE TABLE](#) might be fastest.
- Because the primary key is integral to the storage layout of each InnoDB table,
 - and changing the definition of the primary key involves reorganizing the whole table,
 - **always** set up the primary key as part of the [CREATE TABLE](#) statement,
 - and plan ahead so that you do **not** need to **ALTER** or **DROP** the primary key afterward.

- Optimizing InnoDB Disk I/O

- If you follow best practices for database design and tuning techniques for SQL operations, but your database is still slow due to heavy disk I/O activity, consider these disk I/O optimizations.
- If the Unix top tool or the Windows Task Manager shows that the CPU usage percentage with your workload is less than 70%, your workload is probably disk-bound.
- Increase buffer pool size
 - When table data is cached in the InnoDB buffer pool, it can be accessed repeatedly by queries without requiring any disk I/O. Specify the size of the buffer pool with the [innodb buffer pool size](#) option.
- Adjust the flush method
 - If database write performance is an issue, conduct benchmarks with the [innodb flush method](#) parameter set to **O_DSYNC**.
- Configure an **fsync** threshold
 - You can use the [innodb fsync threshold](#) variable to define a threshold value, in bytes.
 - Specifying a threshold to force smaller, periodic flushes may be beneficial in cases where multiple MySQL instances use the same storage devices.

- Optimizing InnoDB Disk I/O
 - Use a **noop** or **deadline** I/O scheduler with native AIO on Linux
 - InnoDB uses the asynchronous I/O subsystem (native AIO) on Linux to perform read-ahead and write requests for data file pages. This behavior is controlled by the [innodb use native aio](#) configuration option, which is enabled by default.
 - Use direct I/O on Solaris 10 for x86_64 architecture
 - To apply direct I/O only to InnoDB file operations rather than the whole file system, set [innodb flush method = 0 DIRECT](#).
 - Use raw storage for data and log files with Solaris 2.6 or later
 - Users of the Veritas file system VxFS should use the **convosync=direct** mount option.
 - Use **additional** storage devices
 - Additional storage devices could be used to set up a RAID configuration.
 - Alternatively, InnoDB tablespace data files and log files can be placed on **different physical disks**.

- Optimizing InnoDB Disk I/O
 - Consider **non-rotational** storage
 - Non-rotational storage generally provides better performance for **random I/O operations**; and rotational storage for sequential I/O operations. When distributing data and log files across rotational and non-rotational storage devices, consider the type of I/O operations that are predominantly performed on each file.
 - Random I/O-oriented files typically include [file-per-table](#) and [general tablespace](#) data files, [undo tablespace](#) files, and [temporary tablespace](#) files.
 - Sequential I/O-oriented files include InnoDB [system tablespace](#) files (due to [doublewrite buffering](#) prior to MySQL 8.0.20 and [change buffering](#)), doublewrite files introduced in MySQL 8.0.20, and log files such as [binary log](#) files and [redo log](#) files.
 - Increase I/O capacity to avoid backlogs
 - If throughput **drops periodically** because of InnoDB [checkpoint](#) operations, consider increasing the value of the [innodb io capacity](#) configuration option.
 - Higher values cause more frequent [flushing](#), avoiding the backlog of work that can cause dips in throughput.

- Optimizing InnoDB Disk I/O

- Lower I/O capacity if flushing does not fall behind

- In a typical scenario where you could **lower** the option value, you might see a combination like this in the output from [SHOW ENGINE INNODB STATUS](#):

- History list length low, below a few thousand.
 - Insert buffer merges close to rows inserted.
 - Modified pages in buffer pool consistently well below [innodb max dirty pages pct](#) of the buffer pool.
 - Log sequence number - Last checkpoint is at less than 7/8 or ideally less than 6/8 of the total size of the InnoDB [log files](#).

- Store **system tablespace** files on Fusion-io devices

- You can take advantage of a **doublewrite** buffer-related I/O optimization by storing the files that contain the doublewrite storage area on Fusion-io devices that support atomic writes.

- Disable logging of compressed pages

- If you are certain that the zlib version is not subject to change, disable [innodb log compressed pages](#) to reduce redo log generation for workloads that modify compressed data.

- Optimizing InnoDB Configuration Variables
 - Controlling the types of data change operations for which InnoDB buffers the changed data, to **avoid frequent small disk writes**.
 - Turning the **adaptive hash indexing feature** on and off using the [innodb adaptive hash index](#) option.
 - Setting a limit on **the number of concurrent threads** that InnoDB processes, if context switching is a bottleneck.
 - Controlling **the amount of prefetching** that InnoDB does with its read-ahead operations.
 - Increasing **the number of background threads** for read or write operations, if you have a high-end I/O subsystem that is not fully utilized by the default values
 - Controlling **how much I/O** InnoDB performs in the background.
 - Controlling the algorithm that determines when InnoDB performs **certain types of background writes**.

- Optimizing InnoDB Configuration Variables
 - Taking advantage of multicore processors and their cache memory configuration, to **minimize delays in context switching**.
 - Preventing **one-time operations** such as table scans from interfering with the frequently accessed data stored in the InnoDB buffer cache.
 - Adjusting log files to **a size that makes sense for reliability and crash recovery**.
 - Configuring **the size and number of instances** for the InnoDB **buffer pool**, especially important for systems with multi-gigabyte buffer pools.
 - Increasing **the maximum number of concurrent transactions**, which dramatically improves scalability for the busiest databases.
 - Moving **purge operations** (a type of garbage collection) into a **background** thread.
 - Reducing the amount of switching that InnoDB does **between concurrent threads**, so that SQL operations on a busy server do not queue up and form a “traffic jam”.

- Optimizing InnoDB for Systems with Many Tables
 - If you have configured [non-persistent optimizer statistics](#) (a non-default configuration),
 - InnoDB computes index [cardinality](#) values for a table **the first time that table is accessed after startup**, instead of storing such values in the table.
 - This step can take significant time on systems that partition the data into many tables.
 - Since this overhead only applies to the initial table open operation, to “warm up” a table for later use, access it immediately after startup by issuing a statement such as **SELECT 1 FROM *tbl_name* LIMIT 1**.
 - Optimizer statistics are persisted to disk by default, enabled by the [innodb_stats_persistent](#) configuration option.

- Optimizing MyISAM Queries

- To help MySQL better optimize queries, use [ANALYZE TABLE](#)
 - or run [myisamchk --analyze](#) on a table after it has been loaded with data.
- To sort an index and data according to an index, use [myisamchk --sort-index --sort-records=1](#).
- Try to avoid complex [SELECT](#) queries on MyISAM tables that are updated frequently,
 - to avoid problems with **table locking** that occur due to contention between readers and writers.
- MyISAM supports concurrent inserts:
 - If a table has no free blocks in the middle of the data file, you can [INSERT](#) new rows into it at the same time that other threads are reading from the table. If it is important to be able to do this, consider using the table in ways that avoid deleting rows.
- For MyISAM tables that **change frequently**,
 - try to **avoid all variable-length columns** ([VARCHAR](#), [BLOB](#), and [TEXT](#)). The table uses dynamic row format if it includes even a single variable-length column.

- Optimizing MyISAM Queries

- It is normally **not useful** to split a table into different tables just because the rows become large.
 - In accessing a row, the biggest performance hit is the disk seek needed to find the first byte of the row. After finding the data, most modern disks can read the entire row fast enough for most applications.
- Use **ALTER TABLE ... ORDER BY *expr1*, *expr2***, ... if you usually retrieve rows in *expr1*, *expr2*, ... order.
 - By using this option after extensive changes to the table, you may be able to get higher performance.
- If you often need to calculate results such as counts based on information from a lot of rows, it may be preferable to introduce a new table and update the counter in real time. An update of the following form is very fast:
UPDATE *tbl_name* SET *count_col*=*count_col*+1 WHERE *key_col*=*constant*;
- Use **OPTIMIZE TABLE** periodically to avoid fragmentation with dynamic-format MyISAM tables.
- Declaring a MyISAM table with the **DELAY_KEY_WRITE=1** table option makes index updates faster because they are not flushed to disk until the table is closed.

- Bulk Data Loading for MyISAM Tables

- For a MyISAM table, you can use **concurrent inserts** to add rows at the same time that [SELECT](#) statements are running, if there are no deleted rows in middle of the data file.
- With some extra work, it is possible to make [LOAD DATA](#) run even faster for a MyISAM table when the table has **many indexes**. Use the following procedure:
 - Execute a [FLUSH TABLES](#) statement or a [mysqladmin flush-tables](#) command.
 - Use [myisamchk --keys-used=0 -rq /path/to/db/tbl name](#) to remove all use of indexes for the table.
 - Insert data into the table with [LOAD DATA](#). This does not update any indexes and therefore is very fast.
 - If you intend only to read from the table in the future, use [myisampack](#) to compress it.
 - Re-create the indexes with [myisamchk -rq /path/to/db/tbl name](#). This creates the index tree in memory before writing it to disk, which is much faster than updating the index during [LOAD DATA](#) because it avoids lots of disk seeks. The resulting index tree is also perfectly balanced.
 - Execute a [FLUSH TABLES](#) statement or a [mysqladmin flush-tables](#) command.

- Bulk Data Loading for MyISAM Tables
 - [LOAD DATA](#) performs the preceding optimization automatically if the MyISAM table into which you insert data is empty.
 - The main difference between automatic optimization and using the procedure explicitly is that you can let [myisamchk](#) allocate much more temporary memory for the index creation than you might want the server to allocate for index re-creation when it executes the [LOAD DATA](#) statement.
 - You can also **disable** or **enable** the nonunique indexes for a MyISAM table by using the following statements rather than [myisamchk](#). If you use these statements, you can skip the [FLUSH TABLES](#) operations:
ALTER TABLE *tbl_name* DISABLE KEYS;
ALTER TABLE *tbl_name* ENABLE KEYS;

- Bulk Data Loading for MyISAM Tables

- To speed up [INSERT](#) operations that are performed with multiple statements for **nontransactional** tables, lock your tables:

LOCK TABLES a WRITE;

INSERT INTO a VALUES (1,23),(2,34),(4,33);

INSERT INTO a VALUES (8,26),(6,29);

...

UNLOCK TABLES;

- This benefits performance because the index buffer is flushed to disk only once, after all [INSERT](#) statements have completed.
- Locking also **lowers the total time for multiple-connection tests**, although the maximum wait time for individual connections might go up because they wait for locks.
- Suppose that five clients attempt to perform inserts simultaneously as follows:
 - Connection 1 does 1000 inserts**
 - Connections 2, 3, and 4 do 1 insert**
 - Connection 5 does 1000 inserts**
- If you do not use locking, connections 2, 3, and 4 finish before 1 and 5.
- If you use locking, connections 2, 3, and 4 probably do not finish before 1 or 5, but the total time should be about 40% faster.

- Optimizing REPAIR TABLE Statements
 - [REPAIR TABLE](#) for MyISAM tables is similar to using [myisamchk](#) for repair operations
 - Suppose that a [myisamchk](#) table-repair operation is done using the following options to set its memory-allocation variables:
 - key_buffer_size=128M --myisam_sort_buffer_size=256M
 - read_buffer_size=64M --write_buffer_size=64M
 - Some of those [myisamchk](#) variables correspond to server system variables:

myisamchk Variable	System Variable
key_buffer_size	key_buffer_size
myisam_sort_buffer_size	myisam_sort_buffer_size
read_buffer_size	read_buffer_size
write_buffer_size	none

- Consider using MEMORY tables for
 - **noncritical data** that is accessed often, and is **read-only** or **rarely updated**.
- For best performance with MEMORY tables,
 - examine the kinds of queries against each table, and specify the type to use for each associated index, either **a B-tree index** or a **hash index**.
 - On the **CREATE INDEX** statement, use the clause **USING BTREE** or **USING HASH**.
 - B-tree indexes are fast for queries that do greater-than or less-than comparisons through operators such as > or BETWEEN.
 - Hash indexes are only fast for queries that look up single values through the = operator, or a restricted set of values through the IN operator.

- InnoDB Buffer Pool Optimization
 - [InnoDB](#) maintains a storage area called the [buffer pool](#) for caching data and indexes in memory.
 - [Configuring InnoDB Buffer Pool Prefetching \(Read-Ahead\)](#)
 - [Configuring Buffer Pool Flushing](#)
 - [Making the Buffer Pool Scan Resistant](#)
 - [Configuring Multiple Buffer Pool Instances](#)
 - [Saving and Restoring the Buffer Pool State](#)
 - [Configuring InnoDB Buffer Pool Size](#)

- Configuring InnoDB Buffer Pool Size

- When increasing or decreasing [innodb buffer pool size](#), the operation is performed in chunks.
 - Chunk size is defined by the [innodb buffer pool chunk size](#) configuration option, which has a default of **128M**. Buffer pool size must always be equal to or a multiple of [innodb buffer pool chunk size](#) * [innodb buffer pool instances](#).

```
shell> mysql --innodb-buffer-pool-size=8G --innodb-buffer-pool-instances=16
```

```
mysql> SELECT @@innodb_buffer_pool_size/1024/1024/1024;
```

```
+-----+
| @@innodb_buffer_pool_size/1024/1024/1024 |
+-----+
|                                     8.000000000000000 |
+-----+
```

```
shell> mysql --innodb-buffer-pool-size=9G --innodb-buffer-pool-instances=16
```

```
mysql> SELECT @@innodb_buffer_pool_size/1024/1024/1024;
```

```
+-----+
| @@innodb_buffer_pool_size/1024/1024/1024 |
+-----+
|                                     10.000000000000000 |
+-----+
```

- Configuring Multiple Buffer Pool Instances
 - For systems with buffer pools in the **multi-gigabyte** range, dividing the buffer pool into **separate instances can improve concurrency**, by reducing contention as different threads read and write to cached pages.
 - Multiple buffer pool instances are configured using the [innodb buffer pool instances](#) configuration option, and you might also adjust the [innodb buffer pool size](#) value.
 - When the InnoDB buffer pool is large, many data requests can be satisfied by retrieving from memory.
 - To enable multiple buffer pool instances, set the **innodb_buffer_pool_instances** configuration option to a value **greater than 1 (the default) up to 64 (the maximum)**.
 - This option takes effect only when you set **innodb_buffer_pool_size** to a size of **1GB or more**. The total size you specify is divided among all the buffer pools. For best efficiency, specify a combination of [innodb buffer pool instances](#) and [innodb buffer pool size](#) so that each buffer pool instance is **at least 1GB**.

- Making the Buffer Pool Scan Resistant
 - Rather than using a strict [LRU](#) algorithm,
 - InnoDB uses a technique to minimize the amount of data that is brought into the [buffer pool](#) and never accessed again.
 - The goal is to make sure that frequently accessed (“hot”) pages remain in the buffer pool, even as [read-ahead](#) and [full table scans](#) bring in new blocks that might or might not be accessed afterward.
 - Newly read blocks are inserted into **the middle of the LRU list**.
 - All newly read pages are inserted at a location that by default is **3/8 from the tail of the LRU list**.
 - The pages are **moved to the front of the list** (the most-recently used end) when they are accessed in the buffer pool for the first time.
 - Thus, pages that are never accessed never make it to the front portion of the LRU list, and **“age out” sooner** than with a strict LRU approach.
 - This arrangement divides the LRU list into **two segments**, where the pages downstream of the insertion point are considered “old” and are desirable victims for LRU eviction.

- Configuring InnoDB Buffer Pool Prefetching (Read-Ahead)
 - A [read-ahead](#) request is an I/O request to prefetch multiple pages in the [buffer pool](#) **asynchronously**, in anticipation of **impending need** for these pages.
 - The requests bring in all the pages in one [extent](#). InnoDB uses two read-ahead algorithms to improve I/O performance:
 - **Linear** read-ahead is a technique that predicts what pages might be needed soon based on pages in the buffer pool being accessed sequentially.
 - **Random** read-ahead is a technique that predicts when pages might be needed soon based on pages already in the buffer pool, regardless of the order in which those pages were read.
 - The **SHOW ENGINE INNODB STATUS** command displays statistics to help you evaluate the effectiveness of the read-ahead algorithm.
 - The information can be useful when fine-tuning the [innodb random read ahead](#) setting.

- Configuring Buffer Pool Flushing
 - InnoDB performs certain tasks in the background, including flushing of dirty pages from the buffer pool.
 - **Dirty pages** are those that have been modified but are not yet written to the data files on disk.
 - In MySQL 8.0, buffer pool flushing is performed by page cleaner threads.
 - The number of page cleaner threads is controlled by the [innodb page cleaners](#) variable, which has a default value of **4**.
 - However, if the number of page cleaner threads exceeds the number of buffer pool instances, [innodb page cleaners](#) is automatically set to the same value as [innodb buffer pool instances](#).
 - Buffer pool flushing is initiated when the percentage of dirty pages reaches the low water mark value defined by the [innodb max dirty pages pct lwm](#) variable.
 - The default low water mark is **10%** of buffer pool pages. A [innodb max dirty pages pct lwm](#) value of 0 disables this early flushing behaviour.

- Saving and Restoring the Buffer Pool State
 - To reduce the [warmup](#) period after restarting the server, InnoDB saves a percentage of the most recently used pages for each buffer pool at server shutdown and restores these pages at server startup.
 - The percentage of recently used pages that is stored is defined by the [innodb_buffer_pool_dump_pct](#) configuration option.
 - Operations related to saving and restoring the buffer pool state are described in the following topics:
 - [Configuring the Dump Percentage for Buffer Pool Pages](#)
 - [Saving the Buffer Pool State at Shutdown and Restoring it at Startup](#)
 - [Saving and Restoring the Buffer Pool State Online](#)
 - [Displaying Buffer Pool Dump Progress](#)
 - [Displaying Buffer Pool Load Progress](#)
 - [Aborting a Buffer Pool Load Operation](#)
 - [Monitoring Buffer Pool Load Progress Using Performance Schema](#)

- The MyISAM Key Cache

- To minimize disk I/O, MyISAM storage engine employs a cache mechanism to keep the most frequently accessed table blocks in memory:
 - For **index** blocks, a special structure called the **key cache** (or key buffer) is maintained. The structure contains a number of block buffers where the **most-used index blocks** are placed.
 - For **data** blocks, MySQL uses **no special cache**. Instead it relies on the **native operating system file system cache**.
- To control the size of the key cache, use the [key buffer size](#) system variable.
 - If this variable is set equal to **zero**, no key cache is used. The key cache also is **not** used if the [key buffer size](#) value is too small to allocate the minimal number of block buffers (8).
 - When the key cache is **not** operational, index files are accessed using only the native file system buffering provided by the operating system. (In other words, table index blocks are accessed using the same strategy as that employed for table data blocks.)

- Shared Key Cache Access

- Threads can access key cache buffers **simultaneously**, subject to the following conditions:
 - A buffer that is **not** being updated can be accessed by multiple sessions.
 - A buffer that is being updated causes sessions that need to use it to **wait** until the update is complete.
 - Multiple sessions can initiate requests that result in cache block replacements, as long as they **do not interfere with each other** (that is, as long as they need different index blocks, and thus cause different cache blocks to be replaced).
- Shared access to the key cache enables the server to improve throughput significantly.

- Multiple Key Caches

- To reduce key cache access contention further, MySQL also provides **multiple key caches**.
 - This feature enables you to assign **different table indexes to different key caches**.
- Where there are multiple key caches, the server must know which cache to use when processing queries for a given MyISAM table.
 - By default, all MyISAM table indexes are cached in the **default key cache**.
 - To assign table indexes to a specific key cache, use the [CACHE INDEX](#) statement

```
mysql> CACHE INDEX t1, t2, t3 IN hot_cache;
```

Table	Op	Msg_type	Msg_text
test.t1	assign_to_keycache	status	OK
test.t2	assign_to_keycache	status	OK
test.t3	assign_to_keycache	status	OK

- Multiple Key Caches

- For a busy server, you can use a strategy that involves three key caches:
 - A “**hot**” key cache that takes up **20%** of the space allocated for all key caches. Use this for tables that are heavily used for searches but that are not updated.
 - A “**cold**” key cache that takes up **20%** of the space allocated for all key caches. Use this cache for medium-sized, intensively modified tables, such as temporary tables.
 - A “**warm**” key cache that takes up **60%** of the key cache space. Employ this as the default key cache, to be used by default for all other tables.
- The following example assigns several tables each to **hot_cache** and **cold_cache**:
`CACHE INDEX db1.t1, db1.t2, db2.t3 IN hot_cache`
`CACHE INDEX db1.t4, db2.t5, db2.t6 IN cold_cache`

- Midpoint Insertion Strategy

- By default, the key cache management system uses a simple **LRU** strategy for choosing key cache blocks to be evicted, but it also supports a more sophisticated method called the **midpoint insertion strategy**.
- When using the midpoint insertion strategy, the LRU chain is divided into two parts: **a hot sublist** and **a warm sublist**.
 - The division point between two parts is **not fixed**, but the key cache management system takes care that the warm part is not “too short,” always containing at least key cache division limit percent of the key cache blocks.
 - key cache division limit is a component of structured key cache variables, so its value is a parameter that can be set per cache.
- The midpoint insertion strategy helps to **improve performance** when execution of a query that requires an index scan effectively pushes out of the cache all the index blocks corresponding to valuable high-level B-tree nodes.

- Index Preloading

- If there are enough blocks in a key cache to hold blocks of an **entire index**,
 - or at least the blocks corresponding to its **nonleaf nodes**, it makes sense to preload the key cache with index blocks before starting to use it.
 - Preloading enables you to put the table index blocks into a key cache buffer in the most efficient way: **by reading the index blocks from disk sequentially**.
- To preload an index into a cache, use the [LOAD INDEX INTO CACHE](#) statement.
- For example, the following statement preloads nodes (index blocks) of indexes of the tables t1 and t2:

```
mysql> LOAD INDEX INTO CACHE t1, t2 IGNORE LEAVES;
```

+-----+-----+-----+-----+			
Table	Op	Msg_type	Msg_text
+-----+-----+-----+-----+			
test.t1	preload_keys	status	OK
test.t2	preload_keys	status	OK
+-----+-----+-----+-----+			

- Key Cache Block Size

- It is possible to specify the **size of the block buffers** for an individual key cache using the [key cache block size](#) variable.
 - This permits tuning of the performance of I/O operations for index files.
- The best performance for I/O operations is achieved when the size of read buffers is **equal to the size of the native operating system I/O buffers**.
 - But setting the size of key nodes equal to the size of the I/O buffer does **not** always ensure the best overall performance.
 - When reading the **big leaf nodes**, the server pulls in a lot of unnecessary data, effectively preventing reading other leaf nodes.
- To control the size of blocks in the .MYI index file of MyISAM tables, use the [--myisam-block-size](#) option at server startup.

- Restructuring a Key Cache

- A key cache can be restructured at any time by updating its parameter values. For example:
`mysql> SET GLOBAL cold_cache.key_buffer_size=4*1024*1024;`
- If you assign to either the [key buffer size](#) or [key cache block size](#) key cache component a value that differs from the component's current value,
 - the server **destroys** the cache's old structure and creates a new one based on the new values.
 - If the cache contains any **dirty** blocks, the server saves them to disk before destroying and re-creating the cache.
 - Restructuring does **not** occur if you change other key cache parameters.
- Restructuring does **not** block queries that need to use indexes assigned to the cache.
 - Instead, the server **directly accesses** the table indexes using **native file system caching**. File system caching is not as efficient as using a key cache, so although queries execute, a slowdown can be anticipated.

- Caching of Prepared Statements and Stored Programs
 - **Prepared statements**, both those processed at the SQL level (using the [PREPARE](#) statement) and those processed using the binary client/server protocol (using the [mysql_stmt_prepare\(\)](#) C API function).
 - The [max_prepared_stmt_count](#) system variable controls the total number of statements the server caches. (The sum of the number of prepared statements across all sessions.)
 - **Stored programs** (stored procedures and functions, triggers, and events). In this case, the server converts and caches the entire program body.
 - The [stored_program_cache](#) system variable indicates the approximate number of stored programs the server caches per session.
 - The server maintains caches for prepared statements and stored programs on a **per-session basis**. Statements cached for one session are **not accessible** to other sessions. When a session ends, the server discards any statements cached for it.

- 请你根据上课内容，针对你在E-BookStore项目中的数据库设计，详细回答下列问题：
 1. 你认为在你的数据库中应该建立什么样的索引？为什么？
 2. 你的数据库中每个表中的字段类型和长度是如何确定的？为什么？
 3. 你认为在我们大二上课时讲解ORM映射的Person例子时，每个用户的邮箱如果只有一个，是否还有必要像上课那样将邮箱专门存储在一张表中，然后通过外键关联？为什么？
 4. 你认为主键使用自增主键和UUID各自的优缺点是什么？
 5. 请你搜索参考文献，总结InnoDB和MyISAM两种存储引擎的主要差异，并详细说明你的E-BookStore项目应该选择哪种存储引擎。
- 请提交包含上述问题答案的文档，文档中附上你更新的数据库的设计方案，包括库结构、表结构和表与表之间的关联
- 评分标准：
 - 上述每个问题 1 分，答案不唯一，只要你的说理合理即可视为正确。



Thank You!