

Architecture of Enterprise Applications 23

Virtualization & Container

Haopeng Chen

***RE**liable, **IN**telligent and **SC**alable Systems Group (**REINS**)*

Shanghai Jiao Tong University

Shanghai, China

<http://reins.se.sjtu.edu.cn/~chenhp>

e-mail: chen-hp@sjtu.edu.cn

- Kubernetes
 - Basic Concepts
 - Minikube
- Objectives
 - 能够根据系统容器化部署的需求, 了解K8S的基本原理

- **What is Kubernetes?**

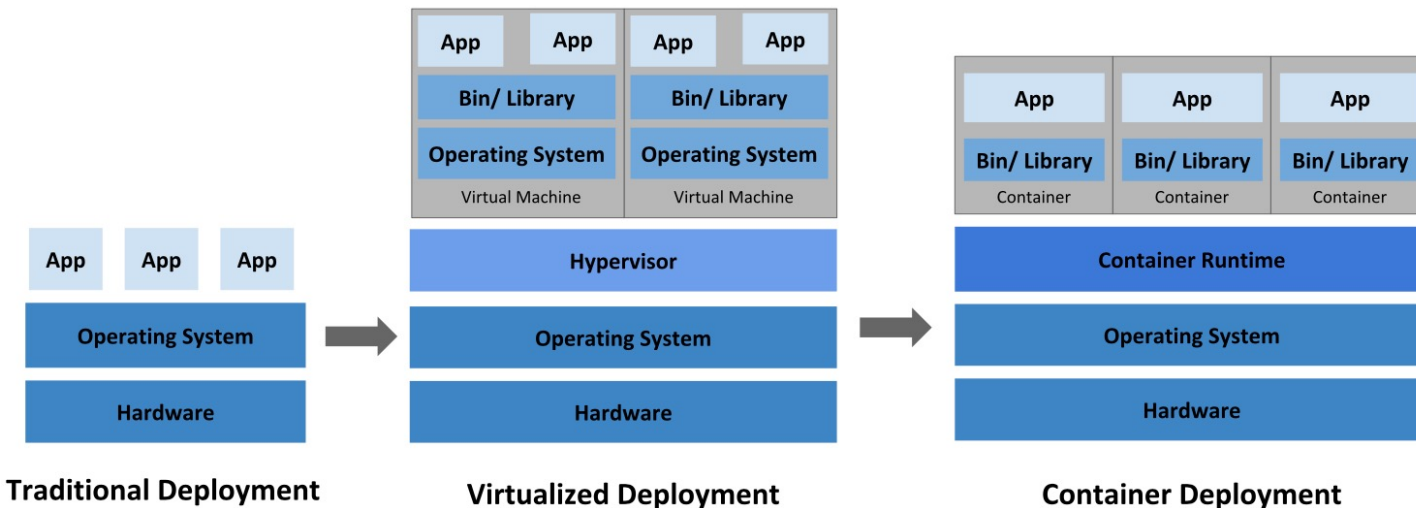
- <https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/>
- Kubernetes is a portable, extensible, open-source platform for **managing containerized workloads and services**, that facilitates both declarative configuration and automation.
- It has a large, rapidly growing ecosystem. Kubernetes services, support, and tools are widely available.
- The name Kubernetes originates from Greek, meaning helmsman or pilot.
- Google open-sourced the Kubernetes project in 2014.
 - Kubernetes combines over 15 years of Google's experience running production workloads at scale with best-of-breed ideas and practices from the community.



kubernetes

• What is Kubernetes?

- Containers are a good way to bundle and run your applications. In a production environment, you need to manage the containers that run the applications and ensure that there is no downtime.
 - For example, if a container goes down, another container needs to start. Wouldn't it be easier if this behavior was handled by a system?

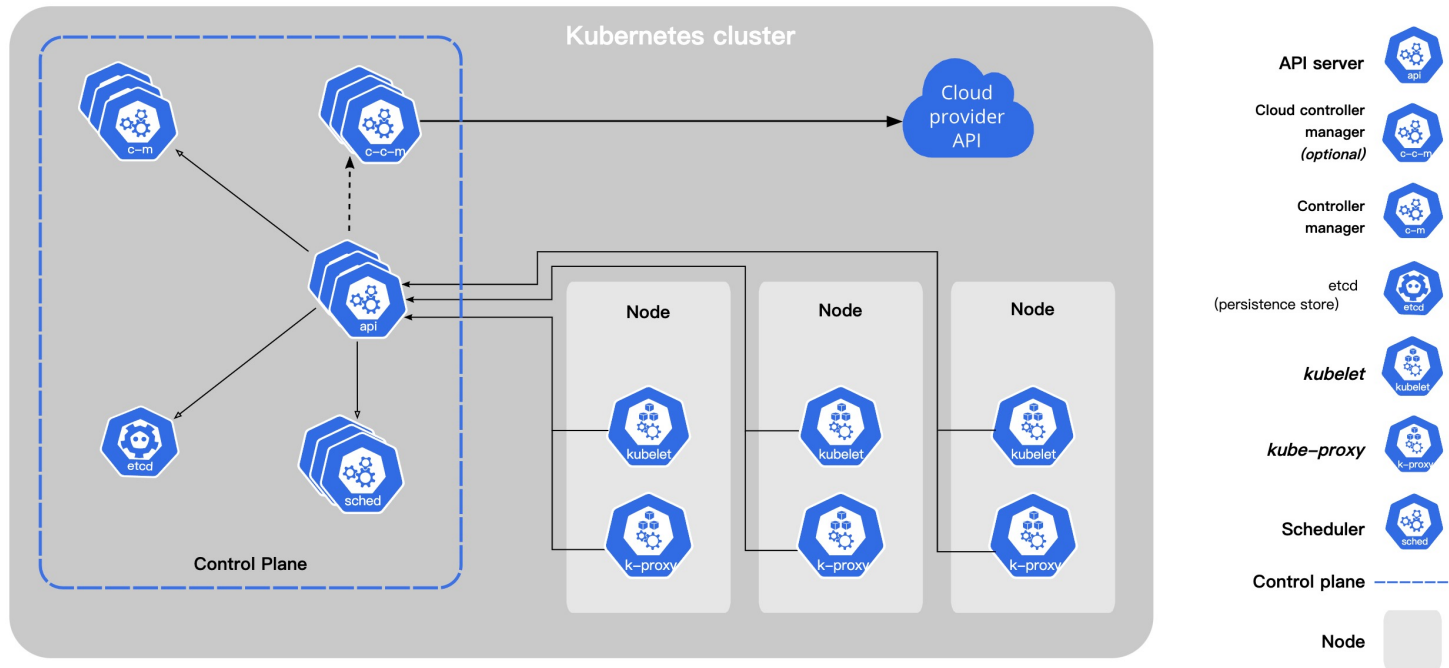


- Kubernetes provides you with:
 - Service discovery and load balancing.
 - Storage orchestration.
 - Automated rollouts and rollbacks.
 - Automatic bin packing.
 - Self-healing .
 - Secret and configuration management.



Kubernetes Component

- When you deploy Kubernetes, you get a **cluster**.
 - A Kubernetes cluster consists of a set of worker machines, called nodes, that **run containerized applications**.
 - Every cluster has **at least one worker node**.



- The control plane's components
 - make **global decisions** about the cluster (for example, scheduling), as well as detecting and responding to cluster events (for example, starting up a new [pod](#) when a deployment's replicas field is unsatisfied).
 - Control plane components can be run **on any machine in the cluster**.
 - However, for simplicity, set up scripts typically start all control plane components **on the same machine**, and do **not** run user containers on this machine.
- kube-apiserver
 - The API server is a component of the Kubernetes [control plane](#) that **exposes the Kubernetes API**. The API server is **the front end** for the Kubernetes control plane.
 - The main implementation of a Kubernetes API server is [kube-apiserver](#). kube-apiserver is designed to scale horizontally—that is, it scales by deploying more instances. You can run **several** instances of kube-apiserver and **balance** traffic between those instances.

- etcd
 - **Consistent and highly-available key value store** used as Kubernetes' backing store for all cluster data.
 - If your Kubernetes cluster uses **etcd** as its backing store, make sure you have a [back up](#) plan for those data.
- kube-scheduler
 - Control plane component that watches for **newly created** [Pods](#) with **no** assigned [node](#), and selects a node for them to run on.
 - **Factors** taken into account for scheduling decisions include: individual and collective resource requirements, hardware/software/policy constraints, affinity and anti-affinity specifications, data locality, inter-workload interference, and deadlines.

- kube-controller-manager
 - Control plane component that runs [controller](#) processes.
 - Logically, each [controller](#) is a **separate process**, but to reduce complexity, they are all compiled into a single binary and run in a single process.
 - Some types of these controllers are:
 - Node controller: Responsible for noticing and responding when nodes go down.
 - Job controller: Watches for Job objects that represent one-off tasks, then creates Pods to run those tasks to completion.
 - Endpoints controller: Populates the Endpoints object (that is, joins Services & Pods).
 - Service Account & Token controllers: Create default accounts and API access tokens for new namespaces.

- cloud-controller-manager
 - A Kubernetes [control plane](#) component that **embeds cloud-specific control logic**.
 - The cloud controller manager lets you **link your cluster into your cloud provider's API**, and separates out the components that interact with that cloud platform from components that only interact with your cluster.
 - The cloud-controller-manager only runs controllers that are specific to your cloud provider.
 - If you are running Kubernetes on your own premises, or in a learning environment inside your own PC, the cluster does not have a cloud controller manager.
 - As with the **kube-controller-manager**, the **cloud-controller-manager** combines several logically independent control loops into a single binary that you run as a single process.
 - You can scale horizontally (run more than one copy) to improve performance or to help tolerate failures.
 - The following controllers can have cloud provider dependencies:
 - Node controller: For checking the cloud provider to determine if a node has been deleted **in the cloud** after it stops responding
 - Route controller: For setting up routes **in the underlying cloud infrastructure**
 - Service controller: For creating, updating and deleting **cloud provider load balancers**

- Node components
 - run on **every node**, maintaining running pods and providing the Kubernetes runtime environment.
- kubelet
 - An agent that runs on each node in the cluster.
 - It makes sure that containers are running in a Pod.
 - The **kubelet** takes a set of **PodSpecs** that are provided through various mechanisms and ensures that the containers described in those **PodSpecs** are running and healthy.
 - The kubelet **doesn't** manage containers which were **not** created by Kubernetes.

- kube-proxy
 - kube-proxy is a network proxy that runs **on each node** in your cluster, implementing part of the Kubernetes [Service](#) concept.
 - [kube-proxy](#) maintains **network rules on nodes**. These network rules allow network communication to your Pods from network sessions inside or outside of your cluster.
 - kube-proxy uses the **operating system packet filtering layer** if there is one and it's available. Otherwise, kube-proxy forwards the traffic itself.
- Container runtime
 - The container runtime is the software that is responsible for **running containers**.
 - Kubernetes supports several container runtimes: [Docker](#), [containerd](#), [CRI-O](#), and any implementation of the [Kubernetes CRI \(Container Runtime Interface\)](#).

- A workload is an application running on Kubernetes.
 - Whether your workload is a single component or several that work together, on Kubernetes you run it inside a set of [pods](#).
 - In Kubernetes, a Pod represents a set of running [containers](#) on your cluster.
- Kubernetes pods have a [defined lifecycle](#).
 - For example, once a pod is running in your cluster then a critical fault on the [node](#) where that pod is running means that all the pods on that node fail.
 - Kubernetes treats that level of failure as final: you would need to create a new Pod to recover, even if the node later becomes healthy.
- However, to make life considerably easier, you don't need to manage each Pod directly.
 - Instead, you can use *workload resources* that manage a set of pods on your behalf.
 - These resources configure [controllers](#) that make sure the right number of the right kind of pod are running, to match the state you specified.

- Kubernetes provides several built-in workload resources:
 - [Deployment](#) and [ReplicaSet](#) (replacing the legacy resource [ReplicationController](#)).
 - Deployment is a good fit for managing a stateless application workload on your cluster, where any Pod in the Deployment is interchangeable and can be replaced if needed.
 - [StatefulSet](#) lets you run one or more related Pods that do track state somehow.
 - For example, if your workload records data persistently, you can run a StatefulSet that matches each Pod with a [PersistentVolume](#). Your code, running in the Pods for that StatefulSet, can replicate data to other Pods in the same StatefulSet to improve overall resilience.
 - [DaemonSet](#) defines Pods that provide node-local facilities.
 - These might be fundamental to the operation of your cluster, such as a networking helper tool, or be part of an [add-on](#). Every time you add a node to your cluster that matches the specification in a DaemonSet, the control plane schedules a Pod for that DaemonSet onto the new node.
 - [Job](#) and [CronJob](#) define tasks that run to completion and then stop.
 - Jobs represent one-off tasks, whereas CronJobs recur according to a schedule.

- *Pods* are the smallest deployable units of computing that you can create and manage in Kubernetes.
 - A *Pod* (as in a pod of whales or pea pod) is a group of one or more [containers](#), with shared storage and network resources, and a specification for how to run the containers.
 - A Pod's contents are always co-located and co-scheduled, and run in a shared context.
 - A Pod models an application-specific "logical host": it contains one or more application containers which are relatively tightly coupled.
 - In non-cloud contexts, applications executed on the same physical or virtual machine are analogous to cloud applications executed on the same logical host.
- As well as application containers, a Pod can contain [init containers](#) that run during Pod startup.
 - You can also inject [ephemeral containers](#) for debugging if your cluster offers this.

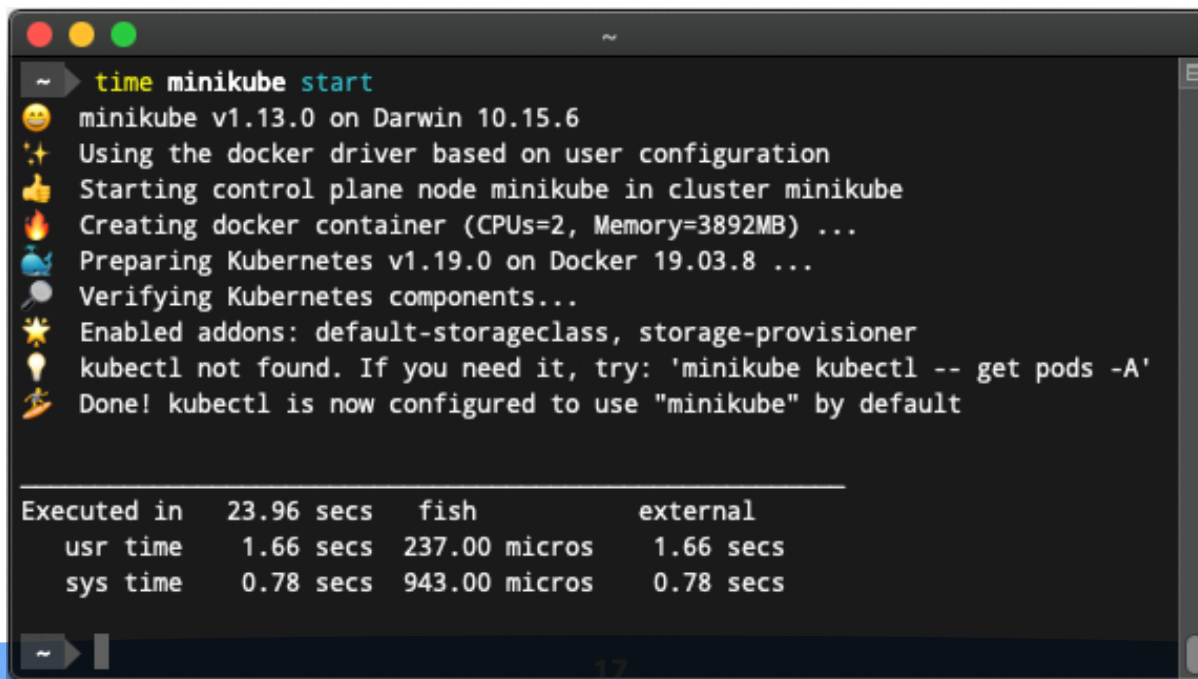
- A *Deployment* provides declarative updates for [Pods](#) and [ReplicaSets](#).
 - You describe a *desired state* in a Deployment, and the Deployment [Controller](#) changes the actual state to the desired state at a controlled rate.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

[controllers/nginx-deployment.yaml](#)

- A Deployment named nginx-deployment is created, indicated by the `.metadata.name` field.
- The Deployment creates three replicated Pods, indicated by the `.spec.replicas` field.
- The `.spec.selector` field defines how the Deployment finds which Pods to manage. In this case, you select a label that is defined in the Pod template (app: nginx).
- The template field contains the following sub-fields:
 - The Pods are labeled `app: nginx` using the `.metadata.labels` field.
 - The Pod template's specification, or `.template.spec` field, indicates that the Pods run one container, nginx, which runs the nginx [Docker Hub](#) image at version 1.14.2.
 - Create one container and name it nginx using the `.spec.template.spec.containers[0].name` field.

- minikube quickly sets up a local Kubernetes cluster on macOS, Linux, and Windows. We proudly focus on helping application developers and new Kubernetes users.
- <https://minikube.sigs.k8s.io/docs/>

A screenshot of a macOS terminal window with a dark background. The window title bar shows red, yellow, and green window control buttons. The prompt is '~'. The command 'time minikube start' has been executed. The output consists of several lines, each preceded by a small emoji icon: a smiley face for the version and OS, a star for the driver, a thumbs up for the control plane, a fire for creating the container, a rocket for preparing Kubernetes, a magnifying glass for verifying components, a star for addons, a lightbulb for the missing kubectl, and a checkmark for completion. Below the output, a table shows the execution time breakdown. The terminal has a scrollbar on the right and a status bar at the bottom showing '~' and '17'.

Operating
system

Linux

macOS

Windows

Architecture

x86-64

ARM64

Release type

Stable

Beta

Installer type

Binary download

Homebrew

To install the latest minikube stable release on x86-64 macOS using binary download:

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-darwin-amd64  
sudo install minikube-darwin-amd64 /usr/local/bin/minikube
```

- Start your cluster
 - From a terminal with administrator access (but not logged in as root), run:
 - `$ minikube start`
- Interact with your cluster
 - If you already have kubectl installed, you can now use it to access your shiny new cluster:
 - `$ kubectl get po -A`
 - Alternatively, minikube can download the appropriate version of **kubectl** and you should be able to use it like this:
 - `$ minikube kubectl -- get po -A`
 - Initially, some services such as the **storage-provisioner**, may not yet be in a Running state.
 - For additional insight into your cluster state, minikube bundles the Kubernetes Dashboard, allowing you to get easily acclimated to your new environment:
 - `$ minikube dashboard`

- Create a sample deployment and expose it on port 8080:
 - `kubectl create deployment hello-minikube --image=k8s.gcr.io/echoserver:1.4`
 - `kubectl expose deployment hello-minikube --type=NodePort --port=8080`
- It may take a moment, but your deployment will soon show up when you run:
 - `kubectl get services hello-minikube`
- The easiest way to access this service is to let minikube launch a web browser for you:
 - `minikube service hello-minikube`
- Alternatively, use kubectl to forward the port:
 - `kubectl port-forward service/hello-minikube 7080:8080`
- Tada! Your application is now available at <http://localhost:7080/>.

- To access a LoadBalancer deployment, use the “minikube tunnel” command. Here is an example deployment:
 - `kubectl create deployment balanced --image=k8s.gcr.io/echoserver:1.4`
 - `kubectl expose deployment balanced --type=LoadBalancer --port=8080`
- In another window, start the tunnel to create a routable IP for the ‘balanced’ deployment:
 - `minikube tunnel`
- To find the routable IP, run this command and examine the EXTERNAL-IP column:
 - `kubectl get services balanced`
- Your deployment is now available at <EXTERNAL-IP>:8080

- Pause Kubernetes without impacting deployed applications:
 - `minikube pause`
- Unpause a paused instance:
 - `minikube unpause`
- Halt the cluster:
 - `minikube stop`
- Increase the default memory limit (requires a restart):
 - `minikube config set memory 16384`
- Browse the catalog of easily installed Kubernetes services:
 - `minikube addons list`
- Create a second cluster running an older Kubernetes release:
 - `minikube start -p aged --kubernetes-version=v1.16.1`
- Delete all of the minikube clusters:
 - `minikube delete --all`

- Kubernetes Documentation
 - <https://kubernetes.io/docs/home/>
- Kubernetes中文手册
 - <https://www.kubernetes.org.cn/docs>
 - <http://docs.kubernetes.org.cn/227.html>
- 推荐一款Kubernetes神器 “minikube”
 - <https://zhuanlan.zhihu.com/p/112755080>
- External IP is pending——ambassador学习
 - <https://blog.csdn.net/TTT12137/article/details/116989159>



Thank You!